

Machine Learning & The Titanic Data

By Souky, Jin, Julie, Cesur, and Emma
April 18, 2020



Agenda:

1. **Tableau Graphics**
2. **Pandas Charts**
3. **Machine Learning Overview and Findings**
4. **Project Conclusions**




Tableau & Pandas Storytelling


- How many people Survived and how many people died ?
- Survival Analysis by Gender .
- Survival Analysis by Class .
- Survival Analysis by Port Embarkation.
- Survival rate per each factors
- Passenger with which traits is likely to



Survived passengers




Passengers by Gender



Survived: Analysis by Gender




Survival Rate Passenger by Gender & Class



Age Survival




Survived: Age

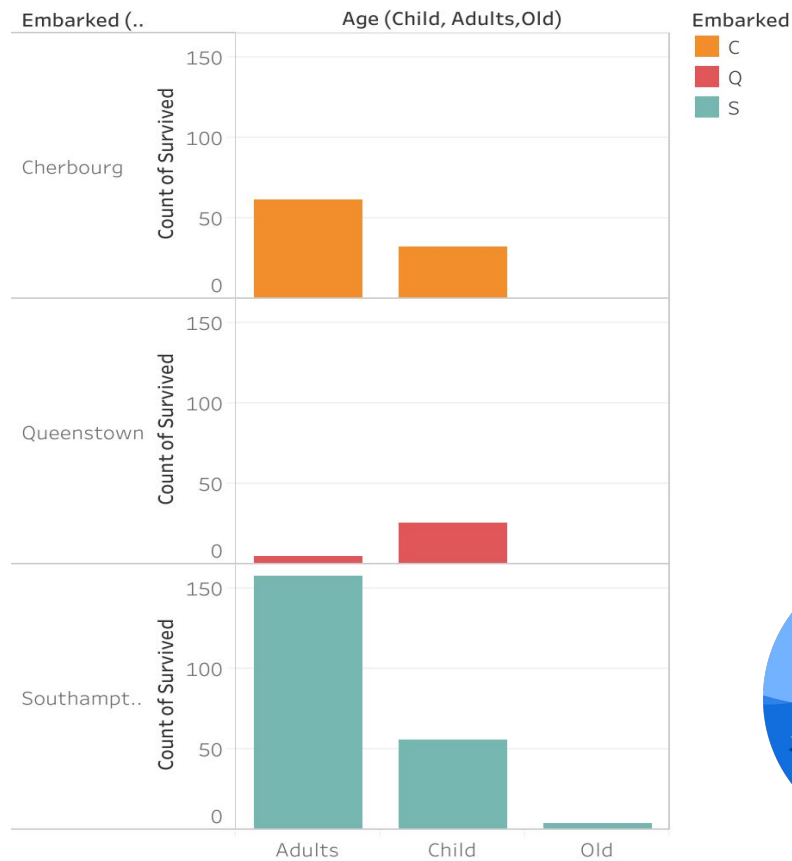


Survived by Class and Embarkation

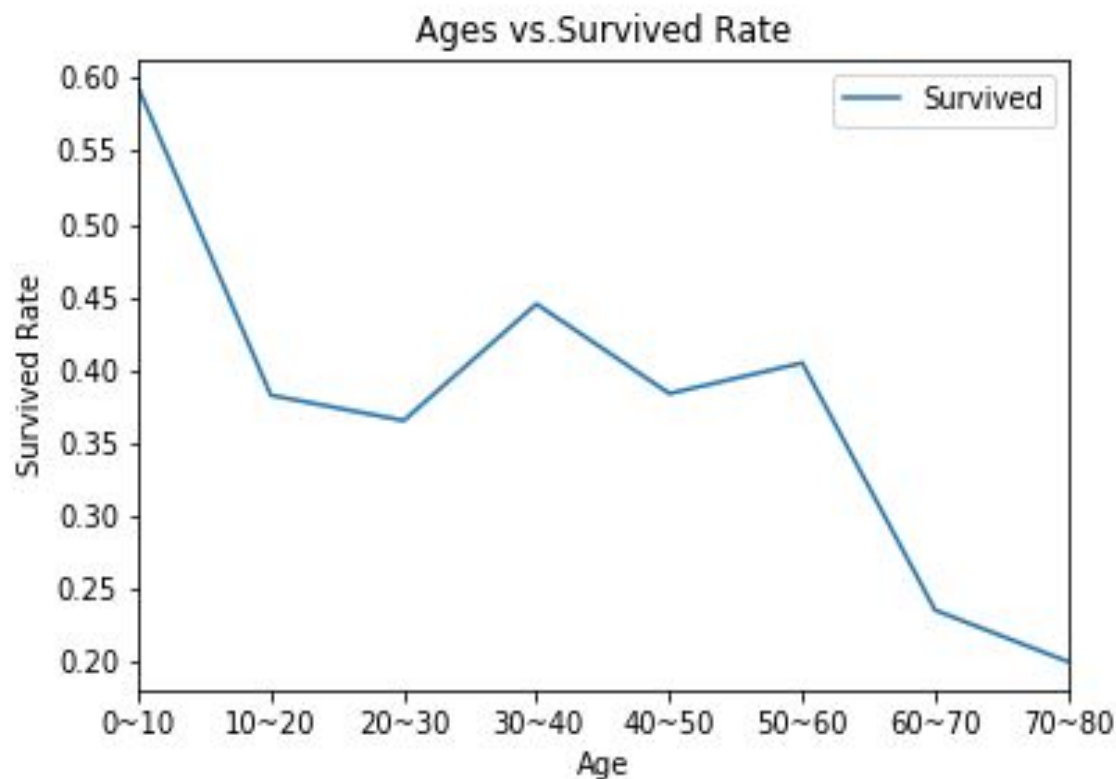
Survived by class





Survived by port of Embarkation




Pandas: Survived rate per factors





Who is more likely to survive?



	Pclass	Sex	Age
Survive%			
100.0	1	female	10~20
100.0	1	female	30~40
100.0	1	female	50~60
100.0	1	female	60~70
100.0	1	male	0~10
100.0	2	female	0~10
100.0	2	female	10~20
100.0	2	male	0~10


	Pclass	Sex	Age
Survive%			
0.0	1	male	60~70
0.0	2	male	20~30
0.0	2	male	50~60
0.0	3	female	40~50
0.0	3	male	50~60
0.0	3	male	60~70
0.0	3	male	70~80

The Conclusion of our Story



Our ML Objective

1. Determine the best machine learning algorithm for our Titanic dataset by utilizing different feature engineering and data cleaning methods
2. Have fun with machine learning
 - a. undergo multiple trial and error tests to experiment with best case scenarios (playing with multiple algorithms and how the code affects accuracy)






ML: Random Forest

```
x_train : input_variables_values_training_datasets
```

```
y_train : target_variables_values_training_datasets
```

								importance	
	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	feature	
0	0	3	0	22.0	1	0	7.2500	Fare	0.297
1	1	1	1	38.0	1	0	71.2833	Age	0.269
2	1	3	1	26.0	0	0	7.9250	Sex	0.268
3	1	1	1	35.0	1	0	53.1000	Pclass	0.084
4	0	3	0	35.0	0	0	8.0500	SibSp	0.045



ML: Confusion Matrix and Comparing Algorithms (Logistic Regression, Decision Tree, and Random Forest)

Estimate Accuracy Scores

Random Forest: 98.2 %

Decision Tree: 98.2 %

Logistic Regression: 79.8 %

CORRECT

INCORRECT

Not Survived, Not Survived

INCORRECT

CORRECT

Survived, Survived

```
# testing the model using confusion matrix
```

```
predictions = cross_val_predict(random_forest, X_train, Y_train)
confusion_matrix(Y_train, predictions)
```

```
array([[475, 74],
       [ 91, 251]])
```

Mean Score: 81.5%

```
predictions = cross_val_predict(decision_tree, X_train, Y_train)
confusion_matrix(Y_train, predictions)
```

```
array([[449, 100],
       [102, 240]])
```


Mean Score: 77.3%

```
predictions = cross_val_predict(logreg, X_train, Y_train)
confusion_matrix(Y_train, predictions)
```

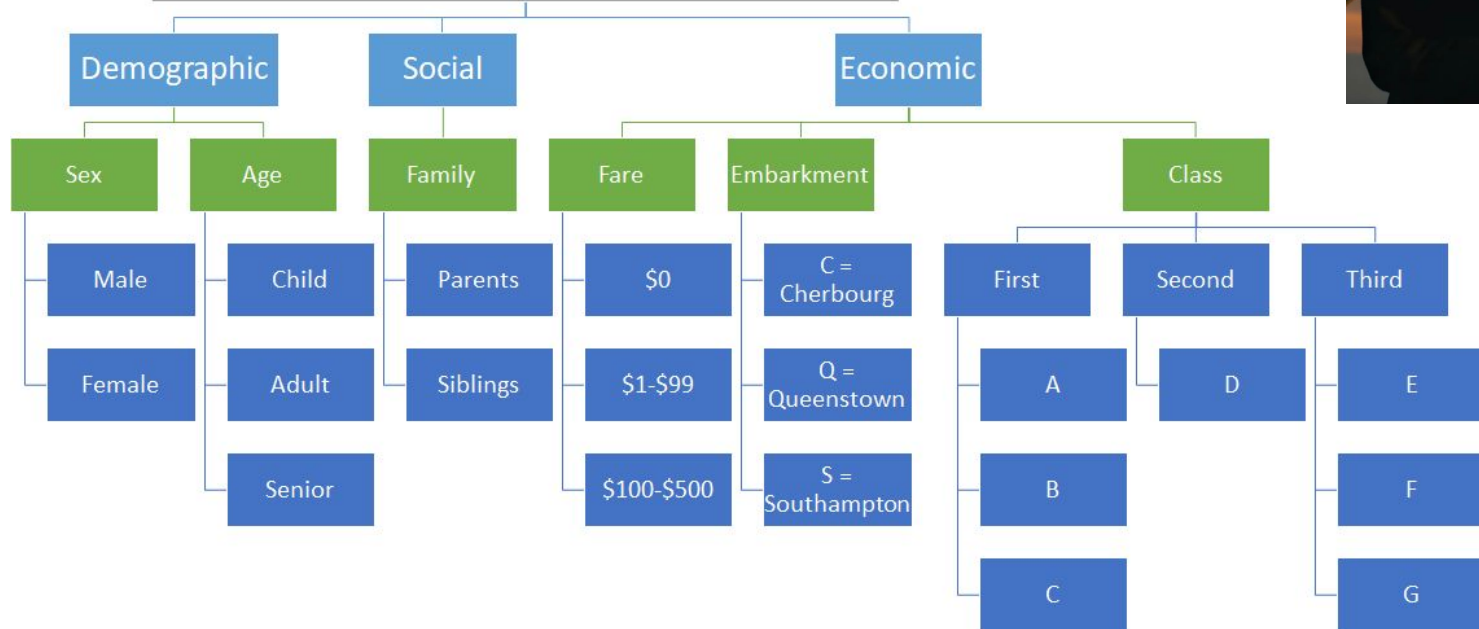
```
array([[463, 86],
       [106, 236]])
```

Mean Score: 78.8%

Passenger Data Analyzing





Attributes of Each Passenger




ML models applied - Decision Tree

- **Decision tree** uses a tree-like model of decisions and their possible consequences, including chance event outcomes, resource costs, and utility.
- Keyword: **Branch**



ML models applied - K Nearest Neighbor




- **K nearest neighbors** is a simple algorithm that stores all available cases and classifies new cases based on a similarity measure.
- Keyword: **Cluster**



<https://www.analyticsvidhya.com/blog/2018/03/introduction-k-neighbours-algorithm-clustering/>



The process of looking for relevant features



	Total	%
Cabin	687	77.1
Age	177	19.9
Embarked	2	0.2
Fare	0	0.0
Ticket	0	0.0



Score and fun discovery

	Model	Score
0	Decision Tree	92.70
1	K Nearest Neighbour	86.76

```
for dataset in data:
    mean = train_df["Age"].mean()
    std = test_df["Age"].std()
    is_null = dataset["Age"].isnull().sum()
    # compute random numbers between the mean, std and is_null
    rand_age = np.random.randint(mean - std, mean + std, size = is_null)
    # fill NaN values in Age column with random values generated
    age_slice = dataset["Age"].copy()
    age_slice[np.isnan(age_slice)] = rand_age
```



Data Preprocessing in Pandas


```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 891 entries, 0 to 890  
Data columns (total 12 columns):  
PassengerId      891 non-null int64  
Survived         891 non-null int64  
Pclass           891 non-null int64  
Name             891 non-null object  
Sex              891 non-null object  
Age             714 non-null float64  
SibSp           891 non-null int64  
Parch           891 non-null int64  
Ticket          891 non-null object  
Fare            891 non-null float64  
Cabin           204 non-null object  
Embarked         889 non-null object  
dtypes: float64(2), int64(5), object(5)  
memory usage: 83.6+ KB
```

```
1 embark=pd.get_dummies(titanic_data["Embarked"],drop_first=True)  
2 embark_test=pd.get_dummies(test_data["Embarked"],drop_first=True)  
3 embark.head()
```

	Q	S
0	0	1
1	0	0
2	0	1
3	0	1
4	0	1



Modelling



```
titanic_data=pd.concat([titanic_data,sex,embark,Pc1],axis=1)
titanic_data.head(2)
```

PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Embarked	male	Q	S	2	3
1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	S	1	0	1	0	1

```
test_data.drop(["Sex","Embarked","Name","Ticket","Pclass"],axis=1, inplace=True)
test_data.head(2)
```

	Survived	Age	SibSp	Parch	Fare	male	Q	S	2	3
0	0	22.0	1	0	7.2500	1	0	1	0	1
1	1	38.0	1	0	71.2833	0	0	0	0	0
2	1	26.0	0	0	7.9250	0	0	1	0	1
3	1	35.0	1	0	53.1000	0	0	1	0	0
4	0	35.0	0	0	8.0500	1	0	1	0	1

Score	Model
80.99	Random Forest
79.42	Naive Bayes
78.29	Decision Tree
70.75	Support Vector Machines
70.31	KNN



Project Final Conclusions

1. How someone manipulates data can produce different machine learning outcomes/accuracy scores
 - a. Dropping null values decreases ML accuracy
 - b. ML does not like non-numeric values either
 - c. A “perfect” ratio of features/variables - can have too many or not enough
2. Different ML algorithms can produce vastly different predictions and accuracy scores even when running through the same data
3. You can guess fairly accurate predictions and make fair hypotheses with Pandas (and other visualization tools) on initial data, even before applying ML
4. The more accurate and robust the dataset, the better ML can “learn” from the data to make better, more sound predictions

Thank you, questions?

